

BROOKE ANDERSON

R FOR ENVIRONMENTAL HEALTH RESEARCH

Contents

1	Prerequisites	5
1.1	Set-up	5
2	Organize	9
2.1	R Projects	9
2.2	Directory organization	10
2.3	Learn more	11
3	Package	13
3.1	Components of R packages	13
3.2	Try it out	15
3.3	Sharing	16
3.4	Learn more	17
4	Track	19
4.1	git	19
4.2	GitHub	21
4.3	Terminal	23
4.4	Try it out	23
4.5	Learn more	24

5 *Collect* 27

5.1 *Data packages* 27

5.2 *Open data APIs* 30

5.3 *Learn more* 31

6 *Bibliography* 33

I

Prerequisites

Based on requests from some of the students for this workshop, I've focused here on a few topics relevant to environmental health research: organizing projects and tracking them with version control, creating your own packages, and collecting and processing large datasets relevant to environmental health research. You can download the slides from the workshop by clicking [here](#).

There are some additional topics in R that would also be useful for environmental health researchers that I won't cover here. I would, however, suggest that you look at the latest on tidyverse functions for cleaning and visualizing data (the `dplyr` and `ggplot2` packages are at the heart of this), new developments on working with geospatial data with the `sf` package, creating interactive graphics with `htmlwidgets`, and creating reports, blogs, and books through the `rmarkdown` framework. These are all areas for you to check out as you continue developing how you use R for research.

Click on the **Next** button (or navigate using the links at the top of the page) to continue.

1.1 Set-up

I am assuming that you already have R and RStudio installed on your computer. You may want to check that you have a recent version of both, and if not, update your version before the workshop. Some of the packages and RStudio tools we'll be using will require newer versions of R and RStudio to work. You can run `sessionInfo()` in R to find out the version of R you have installed. Compare this version to the latest R release version listed at the Comprehensive R Archive Network (CRAN).

To try out the examples, you will also need a bit more set-up:

1. Download git
2. Get a GitHub account
3. Install some R packages
4. Download example R Project

This section will walk you through each step.

1. Download git

In the workshop, you will learn how to use **git**.^[^git] Open-source version control software that can be easily integrated with R projects.] To try the examples, you will need to install git to your computer and make sure that your installation of RStudio can find this software, so you can use git for version control for R Projects. There are more details on this process in the “Track” section.

2. Get a GitHub account

You’ll also learn how to share and collaborate on an R Project using **GitHub**.¹ You will need to get a GitHub account to be able to post repositories on GitHub.

¹ **GitHub**. An online platform for directories tracked with the version control software **git**. This platform has become very popular for sharing code projects, as well as collaborating across a team on developing code and software. Other online git platforms exist and are used by some researchers, including **GitLab**. Once you’ve mastered using GitHub, you should be able to easily transfer those skills to other platforms like GitLab.

3. Install some R packages

This booklet uses a number of R packages beyond base R. To install all the packages that you’ll need, run the following code in your version of R:

```
install.packages(c("devtools", "drat", "tigris",
  "ggplot2"))
```

Also, you’ll want to install the `hurricaneexposedata` and `hurricaneexposure` packages, which you need to do in a bit of a different way:

```
library(drat)
addRepo("geanders")
install.packages("hurricaneexposedata")
install.packages("hurricaneexposure")
```

4. Download example R Project

I’ve created a repository on GitHub. You can find this example repository by clicking [here](#). On the page that takes you to, click on the “Clone or download” button and then select “Download ZIP”.

This will download a single zipped file to your computer. When you unzip the file, it will be a special type of directory, an R Project directory. To open the R Project and start on the examples, open RStudio, then go to “File” -> “Open Project”. A pop-up window will open to let you navigate through your files and find an R Project to open. Navigate to the directory you downloaded, which should be called “columbia_env_health_examples” and doubleclick on the file in this directory called “columbia_env_health_examples.Rproj”.

This will open the project. In the “Files” pane of RStudio, you should see some subdirectories for “R” and “data”. These have the example R code and data, respectively, for you to try the examples in this booklet. The code in each

of the R files should run independently, including the code to load all required packages. Figure 1.1 shows what this package should look like once you've downloaded and opened it, as well as some of the files in the project's "R" subdirectory.

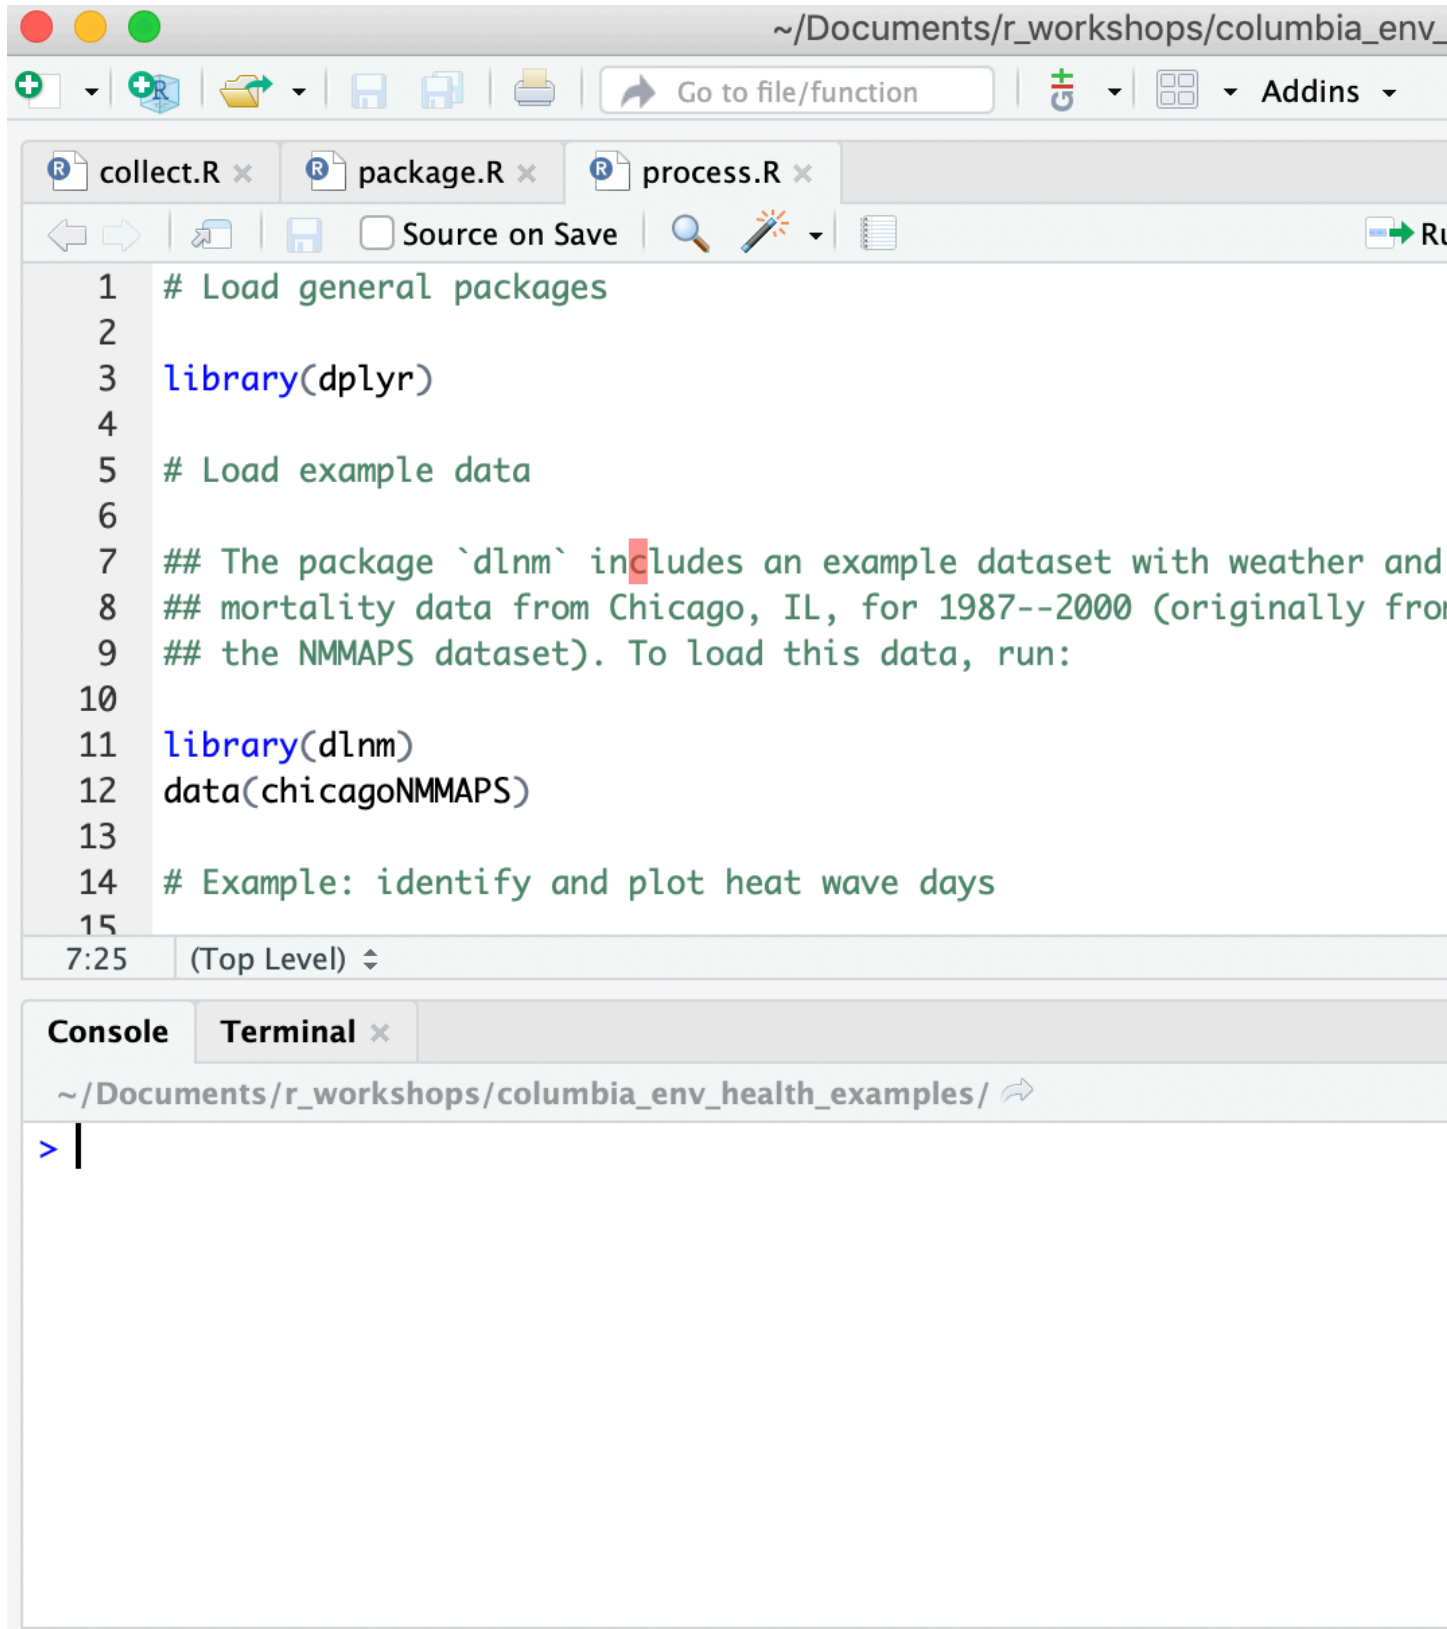


Figure 1.1: What the example R Project for this booklet should look like once you've downloaded and opened it.

2

Organize

If you are using R to write for larger projects, including research for academic papers and theses, you should start putting some thought into how you organize your research files, including raw data, cleaned data, coding scripts for analysis, and output like paper drafts and figures. So far, you may have run much of your analysis within a single R script or R Markdown file. Often, any associated data are within the same working directory as your script or R Markdown file, but the files for one project are not separated from files for other projects. As you move to larger projects, this kind of set-up won't work as well. Instead, you'll want to start keeping all materials for a project in a single and exclusive directory.

2.1 R Projects

RStudio allows you to create “Projects”, which help with this organization. An R Project is very simple—it's just a file directory, with an extra subdirectory added to the directory with settings to remind RStudio about for the project. This directory is saved as a dot file, so you probably won't be able to see it listed if you look at the directory contents using your computer's file finder. If you'd like to see the listing (or delete it by hand, although you likely won't ever need to do this), check the settings for your file finder, and see if you have the option to show all files. Alternatively, you can use the “list” command with the “all” option (`ls -a`) at a Bash shell to view all files and subdirectories in a directory.

Advantages of setting a directory to be an R Project are:

- Automatically uses the directory as your current working directory when you open the project.
- Coordinates well with git version control and GitHub repository system.
- Opens a “Files” window for navigating project files in an RStudio pane when you open the project.

You can either create an R Project as a new directory or convert an existing directory into an R Project. To do either, in RStudio go to the “File” menu and select “New Project”. You’ll then have the option to either create a new directory that’s an R Project or to search through your computer’s files to find an existing directory to make into an R Project.

Once you’ve created an R Project, you can open it in RStudio; opening an R Project will set the project’s directory as your working directory, opening a “Files” window with all the subdirectories and files in the project directory and allowing you to run code with relative filenames from the project’s directory. If you share the R Project with someone else,¹ he or she will also be able to open the R project using RStudio.

One benefit of R Projects is that it is very easy to initialize them as git repositories. A later section (“Track”) will go over how to initialize and use git version control for R Projects. You will also definitely want to use an R Project for any R package you write, as this will introduce a lot of functionality that “plays well” with the `devtools` package to make it easier to write, build, and publish an R package.

¹ One way to do this would be to zip the directory into a single file and share it by email. Another is to use git version control, post the directory to GitHub, and share the directory from there.

2.2 Directory organization

My desk at work is very messy, with lots of paper printouts piled up. But I do keep my project directories very tidy, and I strongly recommend the practice.

This goes beyond “well-organized”, which we just covered (putting all project files in one directory, using subdirectories to divide up files in a project, using consistent names for project file directories, etc.). Keeping a project directory “tidy” means having **only one** version of each file. Often, as you develop a project, especially when collaborators are involved, you can end up with many versions of a file. For example, you may have the draft of a journal article’s text saved in some versions with the file name reflecting the date of the draft (“draft_may_12.docx”), some versions that include the initials of people who reviewed it (“paper_draft_ba_rp_mb2.docx”), and so on.

This type of organization—having multiple versions of project files, with the file names meant to help you keep track of them—results in very cluttered and hard-to-manage project directories. Instead, at any given moment, each of your project directories should have only one version of each project file. You certainly won’t want to lose information from edits and changes to the files along the way, so it’s smart to use some type of **version control system**[**version control system**. Software for tracking changes to software projects. Some researchers who use a lot of coding also use this software to track their projects. One well-known example is **git**, but there are other version control systems as well.] on each project directory. This will allow you to track the changes you’ve made to each file and to go back and revisit the file at any moment in the project’s history. A later section of this booklet (“Track”) will describe how you can use `git` for version control for R projects.

2.2.1 *Avoiding repetition*

One key to efficient organization is to **avoid repetition**.² In practice, it often happens that you're using the same code across several projects. For example, say you have some code that calculates the apparent temperature from air temperature and some measure of dewpoint temperature. If you have organized your project files to have one directory per paper, with all the associated code for a paper within the project's directory, then you may find you're often copying and pasting the code to calculate apparent temperature into different directories.

² In programming, you'll often hear this advice as "Do Not Repeat Yourself".

This situation makes for a tricky balance—you want to organize your files and have a separate directory for each project, but cutting and pasting code can be a recipe for disaster. Each time you move the code, there's a chance for an error to slip in. Also, what if you want to make a change in the code? Say you hear about a better algorithm for calculating apparent temperature? You will either need to go through all of your projects that use that code and change the code everywhere, or you will have to settle for different projects using different algorithms.

This means it's time to start thinking about writing your own R package. You do not have to publish every R package you write—it's fine to just use it yourself and not share it more widely. Regardless, a package is the right place to store related code that you use often, as well as documentation (and possibly tests) to go with that code. A later section of this booklet ("Package") will go over a bit about how to write your own R packages, as well as references for learning more about package development.

2.3 *Learn more*

Reproducible Research with R and R Studio is an excellent book with advice on improving the reproducibility of projects using R, including academic research projects.

A few good articles have come out recently that describe project organization within different scientific fields (e.g., biology, archaeology, ecology). It's worth searching Google Scholar for "reproducible research in R" and your own or a closely related research field.

3

Package

As with many other things in R, the threshold of complexity for writing your own package has recently lowered dramatically. If you are writing some of your own functions in R to use for your research, and you’ve become comfortable with writing functions, you should try putting them in an R package. You don’t have to share this package through publishing it on CRAN or another repository—instead, it’s fine to start by just writing packages for your own use, or for your research group’s private use. However, once you start writing packages, you will see how straightforward it is, and you may want to take the extra steps to prepare your package for CRAN and submit it.

An R package is just “a directory of files which extend R” (<https://cran.r-project.org/doc/manuals/r-release/R-exts.html>). You should consider writing your own package anytime you find yourself cutting and pasting functions to different directories on your computer, so you can use them for different projects. The RStudio team has created a number of tools to make this process easier, including some functionality in the RStudio IDE as well as the `devtools` package (Wickham et al., 2019).

3.1 Components of R packages

3.1.1 R code

Without R code, there’s not much point to an R package.¹ The R function code will go in the package’s “R” subdirectory. You can use as many or as few “.R” files in this section as you want to organize your code. Define your functions in this section as you would if you were creating functions in an R script. As you become more advanced in writing packages, you may want to start adapting your function code a bit. For example, the `package::function` notation can help keep clear which package you mean for a function to call a function from, and so the behavior will be stable regardless of the order that a user loaded packages. Also, if you intend for your functions to be used inside tidyverse pipelines, you’ll want to take a look at the developing Tidy evaluation book.

¹ The exception? A data-only package, where the package’s purpose is to share data rather than R functions. We’ll look at an example in the “Collect” section.

Although it is not required for a package you're creating just to use yourself, you should think about adding help documentation to your code very early. You can then access this documentation through helpfiles in R, just like you do with other R packages. Even if you are just creating a package for yourself, you should add documentation for all the “top-level” functions²

Years ago, you would need to write this documentation in separate “.Rd” files and save them in the “man” subdirectory. Now, the easiest thing to do is to use “roxygen” conventions to write the help information right above each function. With this style, you'll use a special type of commenting character (#') at the start of each line with documentation, and there are special codes (e.g., @param when you're defining one of the function's parameters) to annotate this code. This help documentation will then, when you build your package, be converted to the “.Rd” files you used to have to write yourself and saved in the “man” directory. Several of the resources listed in “Learn more” give more details on writing and rendering this help documentation.

3.1.2 DESCRIPTION file

Every package directory must have a DESCRIPTION file, but if you aren't sharing the package, there's no harm in leaving this as the original template file you get when you start a package directory using RStudio (you'll try this out in just a minute). The exception is when your functions rely on functions from other R packages—in this case, you'll want to add in these dependencies in the DESCRIPTION file. That way, when you load your package, all of these will also be loaded.

If you plan to share your package with others, however, and especially if you plan to publish your package in CRAN or another repository, you'll need to do a lot more with the DESCRIPTION file. This file provides the metadata on your package, including the authors, version, description, and dependencies on other packages. The **R Packages** book referenced in the “Learn More” section has extensive advice on editing this file for packages meant for publication.

3.1.3 Other package components

So far, I've described the barebones elements you'll need to have to create a very minimalist package. There are, however, a number of other elements you'll often want to add to the package. As with the R code for package functions and the DESCRIPTION file, there is a specific location and format required for each of these elements within an R package.

Some of the most common components you'll want to add include example datasets (which will go in a data subdirectory), tutorials for using the package (vignettes subdirectory), and unit tests for the functions (tests subdirectory). To help you set up these extra components, you may want to take a look at the use_* family of functions in the usethis (Wickham and Bryan, 2018) package³ For example, if you want to include example data in your package, you

² **top-level functions.** The functions in your package that you expect your users (including yourself) to call directly. Your package may also include some functions that are “helpers”, only used within other functions to help keep the code clean and avoid replication, but never used directly by package users.

³ The use_* family of functions were previously in the devtools package, but they are being deprecated there and moved to the usethis package. The R Packages book was written before they were deprecated in devtools, and so it discusses them as devtools functions. As long as you install and load the usethis package, though, you'll have no problem following along with the examples in that book.

should check out the `use_data` function, while if you want to add unit tests to your package, check out the `use_test` function. When you run one of these, you'll notice that the function *adds files and subdirectories* when needed to your R package directory. This may throw you at first, since most R functions don't make changes to your file directory, but you'll need those added files to add the package components, and it's *much* easier to remember a single function call to run when you need to add one than to remember all the details of which files (and in what structure) need to be added.

3.2 Try it out

I suggest you try out the following steps to experiment with writing your own package. I've included a function for you to build your package around in the "package.R" file of the R project with examples for the workshop.

- Create a new R Project for your R package
- Build the package
- Move a function into the right place in the package
- Edit the DESCRIPTION file as needed
- Add example data into the right place in the package
- Add some unit tests for the function
- Add documentation using roxygen2 notation
- Add a short vignette

I'll walk you through the first three steps, and then I suggest you use Hadley Wickham's **R Packages** book (see the "Learn More" section) to get some tips as you try out the rest on your own.

Create a new R Project for an R package. Open RStudio and go to "File" -> "New Project". Choose "New Directory" and then "R Package". Name your package "convertr". Then click on "Create Project".

The R Project you've just created will open in RStudio. If you look in the "R" subdirectory, you'll see there's one file called "hello.R". There's also a subdirectory called "man" and three files in the top directory called ".Rbuildignore", "DESCRIPTION", and "NAMESPACE".

Build the package. This project opens with a template—a very small but working package that can serve as a skeleton as you adapt the files to create your own package. Before you do anything to change from the template, try building the package using the `build` function from the `devtools` package:

```
library(devtools)
build()
```

You should get some output like this:

```
checking for file '/Users/georgianaanderson/Documents/r_workshops/convertr/DESCRIPTION' ...
preparing 'convertr':
```

```
checking DESCRIPTION meta-information ...
checking for LF line-endings in source and make files and shell scripts
checking for empty or unneeded directories
building 'convertr_0.1.0.tar.gz'
```

Congratulations! You've built (maybe) your first R package. You should now be able to load it like any other R package on computer and run its function.

Try running:

```
library(convertr)
hello()
```

To build the package in a way that you can use it anywhere on your computer, go to the “Build” tab in RStudio and click “Install and Restart”. As long as you don't get warning messages, you should be able to load and use the package from any R session on your computer.

Move a function into the right place in the package.

Now try adapting the template into your own package. Delete the file “hello.R” in the “R” subdirectory. Create a new file in this subdirectory called something like “temp_conversions.R”. Copy the code for the `c_to_f` function, which is in the “package.R” file of the R Project with examples for this workshop. Save the file and rebuild the package, either with the `build` call to build within that project or with the “Install and Restart” button in the “Build” tab to make it available anywhere on your computer.

3.3 Sharing

3.3.1 When to publish your packages

Once you have some functions that you think others might find useful, you should consider creating a package with them and sharing that package with others. Similarly, if you have a dataset that you think others might find useful and that you have permission to share, you should consider creating a data-only package to use to share that data.

3.3.2 Where to publish packages

There are a few ways to publish your R package, and they have different levels of requirements as well as different levels of recognition (including for CVs and progress reports).

First, you can share the “tar.gz” file that's create when you run “Install and Restart” for an R package from RStudio. The other user can install this file locally using `install.packages` with some extra options. For example, if you shared your “convertr.tar.gz” package file with someone else, they could install it on their computer with the call:⁴

⁴ This call assumes that the other person has “convertr.tar.gz” saved in his or her working directory. Otherwise, use a pathname to direct R to where the file is saved on the computer.


```
install.packages("convertr.tar.gz", repos = NULL,
  type = "source")
```

Second, and perhaps even easier (certainly for the other user), you can post your package as a GitHub repository. The other user can then use the `install_github` function from the `devtools` package to install the package. Unless you set the repository to be private, be aware that your package code would be public. This means anyone could install and use it, although if you don't publicize it, it's unlikely that many people will.

If you want to share your package more widely, you should submit it to a repository. The Comprehensive R Archive Network (CRAN) is the most popular for most scientific fields. The `devtools` package has some tools specifically for checking and releasing a package to CRAN. Bioconductor is the top choice for many bioinformatics packages.

Finally, an excellent choice is to submit the package to ROpenSci for peer review. If you are thinking of doing this, check ROpenSci's policies on what types of packages are considered "in scope" for it. If you submit a package here, it will be openly reviewed by two other R programmers, with an open review process. You will then have the chance to revise the package code in response to the reviewers' comments. Accepted packages can also be submitted to CRAN or Bioconductor.

3.3.3 *Extra steps for publishing*

If you want to submit your package to CRAN, you'll need to take a few extra steps to get it ready. These steps aren't necessary for a package you plan to just use privately, but they wouldn't be a bad idea to try to do even in that case. For example, you'll need to make sure that the functions in your package are comprehensively documented. Every function that you expect an end user might use should have a helpfile. In addition to documenting specific functions, you should also consider writing overall tutorials (called **vignettes**) that walk a user through how to use your package in a few example scenarios. You will also need to make sure that your package can run on different operating systems. If your package only has straightforward functions written exclusively in R, with little interaction outside of R on the user's computer, this shouldn't be an issue. However, you should still check before submitting to CRAN. Several of the resources listed in "Learn More" have very detailed information on preparing a package to submit to a public repository.

3.4 *Learn more*

One wonderful resource for writing and publishing R packages is the book *R Packages* by Hadley Wickham (Wickham, 2015). This book is available for a reasonable price in paperback (many Barnes & Nobles carry it in their computer section) and is also available for free online. The author of this book

also created and maintains the `devtools` **package**⁵ The “Package Development” cheatsheet available through RStudio’s cheatsheets complements these resources.

Along with Roger Peng and Sean Kross, I co-teach a Coursera Specialization on Mastering Software Development in R, and we also have an online book on bookdown on the topic. The official CRAN resource for writing packages, “Writing R Extensions”, is available online. However, it gets very technical, and so I recommend starting with other resources and working up to this document.

To learn more about best practices for creating R, especially for publication, check out ROpenSci’s handbook. ROpenSci conducts peer review for R packages, and it’s all conducted openly. You can learn a lot about packages and package development, especially for more complex packages, by following this peer review package. It’s all conducted through the “Issues” in a GitHub repo.

If you end up writing and publishing a lot of R packages through your scientific research, you may find that you need to help your supervisors and peers understand this research project. One recent article discusses how academic research in the data science realm is changing, as well as how department chairs and others can evaluate these new research products, including for promotion and tenure (Waller, 2018).

⁵ `devtools` **package**. An R package with functions that facilitate R package development. These functions are well-documented in the package documentation as well as through the book *R Packages*. This package also includes a function for installing R packages from code posted on GitHub (`install_github`), which is very useful if you want to find and use packages that are not yet available on CRAN and other repositories or if you want to use the latest, “development” version of a package. This package is available on CRAN and so can be installed with `install.packages("devtools")`.

4

Track

Years ago, I tried to learn to use git version control software with R, and it was a total fail. Maybe it's just me, but I found it really hard to wrap my head around the text-dominated, command-line interface classically used for git. However, RStudio now includes tools that provide a GUI-style interface to most of the functionality you'll need from git for R-based projects. I highly recommend trying git by using it through RStudio first, and then once you develop a mental map of what's going on, it's much easier to transfer partially or completely to running git from a shell.

In this part, I'll discuss both git (the software that allows you to track changes to your R projects), as well as GitHub (an online platform for sharing and collaborating on version-controlled projects). I'll also discuss, near the end, how to use either a Bash Shell or the `git2r` package to do some one-off git tasks that can't be done directly through the RStudio GUI-style git interface.

4.1 *git*

Git is a version control system. It saves information about all changes you make on all files in a repository. This allows you to revert back to previous versions and search through the history for all files in the repository. Git is open source. You can download it for different operating systems at <https://git-scm.com/downloads>. You will need git on your computer to use git with RStudio and create local git repositories you can sync with GitHub repositories. If you are using Windows OS, you'll want your git installation to include a "Bash emulation" (i.e., access to a command line that acts like a Bash shell), which I think should come standard with the Windows git installation. On Mac or Linux computers, you should already have access to a Bash shell.

4.1.1 *Configuring git*

Before you use git, you should configure it. For example, you should make sure it has your name and email address. You can configure git with commands at a

Bash shell.¹ For example, I would run the following code at a shell to configure git to have my proper user name and email:

```
git config --global user.name "Brooke Anderson"
git config --global user.email "brooke.anderson@colostate.edu"
```

You'll only need to run this configuration once, until you get a new computer or for some reason need to re-install git. Once you set these configurations on your computer, they should be set on that computer indefinitely.

Next, you need to make sure that RStudio can find the git software you installed. Sometimes, RStudio will automatically find git (once you've installed git) when you start RStudio. However, in some cases, you may need to take some more steps to activate git in RStudio. To do this, go to "RStudio" -> "Preferences" -> "Git/SVN". Choose "Enable version control". If you have just installed git, and have not restarted RStudio, you'll need to do that before RStudio will recognize git. If you do not see "Git" in the box for "Version control system", it means either that you do not have git installed on your computer or that RStudio was unable to find it. If RStudio doesn't find your version of git in the "Git executable" box, browse for it. If you're still having problems, there's a chance that you'll need to mess around some with your PATH variables. You may want to find a computer-savvy friend or IT person to help you if you get to this stage.

4.1.2 Tracking an R project

To track an R package, you'll first need to initialize it as a git repository. Then, as you make changes to the code, you'll want to **commit** those changes as you go, using **commit messages** that are short but clear enough for you to follow them when you review them later.

You can initialize a git repository for a directory that is an R Project directory through R Studio.

1. Open the Project.
2. Go to "Tools" -> "Version Control" -> "Project Setup".
3. In the box for "Version control system", choose "Git".

You only need to do this once per project, but you will need to do it for each project you want to keep under version control.

Once you initialize the project as a git repository, you should have a "Git" window in one of your RStudio panes (top right pane by default). As you make and save changes to files, they will show up in this window for you to commit. When you want git to record changes, you *commit* the files with the changes. Each time you commit, you have to include a short commit message with some information about the changes. You can make commits from a shell. However, it's much easier to just make commits from the RStudio environment. Click on the "Commit" button in the Git window. That will open a separate commit window. In this window, to commit changes:

¹ If you are using Windows, search for "Bash" among your programs. You should be able to find the Bash emulator that came with your git install. If you are using a Mac, you'll want to open "Terminal". If you're using Linux, you almost certainly already know your way around a Bash shell.

1. Click on the files you want to commit to select them.
2. If you'd like, you can use the bottom part of the window to look through the changes you are committing in each file.
3. Write a message in the "Commit message" box. Keep the message to one line in this box if you can. If you need to explain more, write a short one-line message, skip a line, and then write a longer explanation.
4. Click on the "Commit" button on the right.

Once you commit changes to files, they will disappear from the Git window until you make and save more changes in them.

Once you've made some commits, you may want to browse through the changes. On the top left of the Commit window, you can toggle to "History". This window allows you to explore the history of commits for the repository.

4.2 GitHub

GitHub is an online platform for hosting directories under git version control. It allows you to host git repositories online. This allows you to:

- Work collaboratively on a shared repository
- Fork someone else's repository to create your own copy that you can use and change as you want
- Suggest changes to other people's repositories through pull requests

To push local repositories to GitHub and fork other people's repositories, you will need a GitHub account. You can sign up at <https://github.com>. A free account is fine.

The basic unit for working in GitHub is the repository. You can think of a repository as very similar to an R Project—it's a directory of files with some supplemental files saving some additional information about the directory (and I recommend making any directory with R code that you plan to track with git and GitHub as an R project, although it isn't technically required).

While R Projects have this additional information saved as ".RProj" files, git repositories have this information in a directory called ".git". Because this pathname starts with a dot, it won't show up in many of the ways you list files in a directory.²

If you have a local directory that you would like to push to GitHub, these are the steps to do it. First, you need to make sure that the directory is under git version control. See the notes on initializing a repository. Next, you need to create an empty repository on GitHub to sync with your local repository. Do that by:

1. In GitHub, click on the "+" in the upper right corner ("Create new").
2. Choose "Create new repository".

² From a Bash shell, you can see files that start with . by running `ls -la` from within that directory.

3. Give your repository the same name as the local directory you'd like to connect it to. For example, if you want to connect it to a directory called "fars_analysis" on your computer, name the repository "fars_analysis".
4. Leave everything else as-is (unless you'd like to add a short description in the "Description" box). Click on "Create repository" at the bottom of the page.

Before you connect your local repository with the GitHub one, you'll want to change some settings in RStudio so GitHub will recognize that your local repository belongs to you, rather than asking for your password every time.

- In RStudio, go to "RStudio" -> "Preferences" -> "Git / svn". Choose to "Create RSA key".
- Click on "View public key". Copy what shows up.
- Go to your GitHub account and navigate to "Settings". Click on "SSH and GPG keys".
- Click on "New SSH key". Name the key something like "RStudio" (you might want to include the device name if you'll have SSH keys from RStudio on several computers). Paste in your public key in the "Key box".

Now you are ready to connect the two repositories:

1. Open a shell and navigate to the directory you want to push. (You can open a shell from RStudio using the gear button in the Git window.)
2. Add the GitHub repository as a remote branch with the following command (this gives an example for adding a GitHub repository named "ex_repo" in my GitHub account, "geanders"):

```
git remote add origin git@github.com:geanders/ex_repo.git
```

3. Push the contents of the local repository to the GitHub repository.

```
git push -u origin master
```

Once you have linked a local R project with a GitHub repository, you can push and pull commits using the blue down arrow (pull from GitHub) and green up arrow (push to GitHub) in the Git window in RStudio.

Each original GitHub repository (i.e., not a fork of another repository) has a tab for "Issues". This page works like a Discussion Forum. You can create new "Issue" threads to describe and discuss things that you want to change about the repository. Issues can be closed once the problem has been resolved. You can close issues on the "Issue" page with the "Close issue" button.

GitHub helps you work with others on code. There are two main ways you can do this:

- **Collaborating:** Different people have the ability to push and pull directly to and from the same repository. When one person pushes a change to the repository, other collaborators can immediately get the changes by pulling the latest GitHub commits to their local repository.

- **Forking:** Different people have their own GitHub repositories, with each linked to their own local repository. When a person pushes changes to GitHub, it only makes changes to his own repository. The person must issue a pull request to another person's fork of the repository to share the changes.

See the resources in “Learn more” for more on setting up for this kind of collaboration, including how to resolve merge conflicts when working with collaborators.

4.3 Terminal

I find that, for 90% of what I want to do in git, I can do it through the GUI-style interface RStudio provides for git. The few exceptions include:

- Setting a remote and doing the initial push to that remote
- Reverting a commit
- Creating a new branch
- Merging two branches

For these tasks, I usually open up a Bash shell and run a git command from there. I do this rarely enough that I almost always have to look up the exact call to use to do the task. I've included some resources in “Learn more” that are good to have around when you find you need to open a Bash shell for a git-tracked R project.

There's also a package called `git2r` that lets you run any of these git commands from the R command line, so you could also do all these (fairly rare) tasks from the R console without ever opening a Bash Shell, if you master the `git2r` package. See the package's GitHub page and help documentation to learn more about how to use this package as an alternative to running R from the command line.

4.4 Try it out

4.4.1 Track a project with git

- If you do not already have one, sign up for a GitHub account. The free option is fine.
- If you do not already have git installed on your computer, install it: <https://git-scm.com/downloads>
- Restart RStudio. go to “RStudio” -> “Preferences” -> “Git/SVN”. Choose “Enable version control”. If RStudio doesn't find your version of git in the “Git executable” box, browse for it.
- Open a project you'd like to track with git in RStudio. Change your Project settings to initialize git for this project (see the notes above on how to do that).

- Open a shell from R using the gear symbol in the “Git” pane you should now see in RStudio. Configure git from this shell. For example, I would open a shell and run:

```
git config --global user.name "Brooke Anderson"
git config --global user.email "brooke.anderson@colostate.edu"
```

- Go to the “Commit” window. Click on all of the files you see there, and then make an initial commit using “Initial commit” as your commit message.
- Make a change to one of the files in the project and save the file. You should see that file pop up in the “Git” window of RStudio.
- Commit this change using the “Commit” window. After you commit the changes, look at the “History” window to see the history of your commits.

4.4.2 *Link the project with a GitHub repository*

- Login to your GitHub account online. Create an empty GitHub repository for the project (see the notes above). Give it the same name as the name of your R project directory.
- If you do not already have an RSA key, create one in RStudio and add it as an SSH key in your GitHub settings. If you already have a key (you almost certainly know if you do), see if you can copy it and submit it in GitHub. (You will only need to do this step once, not every time you sync a local git repository with a GitHub repository.)
- Set this empty online GitHub repository as the remote branch of your local git repository for the project using a `git remote add origin` call at a Bash shell (see the notes above).
- Push your local repository to the GitHub repository using a `git push` call (see the notes above).
- Go to your GitHub account and make sure the repository was pushed.
- Try making some more changes to your local repository. Commit the changes, then use the green up arrow in the Git window to push the changes to your GitHub repository.

4.5 *Learn more*

Hadley Wickham’s excellent book on R Packages includes a great chapter on Git and GitHub. While the chapter (and book) is focused on R packages specifically, the guidance in this chapter would apply to any R Project directory under version control, whether or not the R Project is for a package. This book is available free online or as a print version through O’Reilly (and carried at many Barnes & Nobles).

If you’re using git a lot for R projects, it’s helpful to have some resources available with more on using git through a Bash Shell. I like the book *Pragmatic Guide to Git* by Travis Swicegood for a quick, short reference and *Git in Practice* by Mike McQuaid if I’m trying to dig a bit deeper and figure out how git

works. I have also heard good things about Pro Git by Scott Chacon and Ben Straub, which is available for free online.

However, these resources are all geared to those using git for its original purpose of tracking changes in software development projects. When you use git to track files for a research project (rather than, for example, for a repository for an R package you're writing), you may have some needs that don't come up under software development. A few articles have come out recently that give advice on using git specifically to track research projects—these provide helpful advice specific to this use of git, which you won't find in books about git, including one by RStudio's Jenny Bryan (Bryan, 2018).

StackOverflow is also invaluable to quickly look up how to do something in git. There are many tasks in git where I never remember the command, but I do remember enough about what the functionality is called to be able to quickly use Google to find a StackOverflow thread that gives me the call. Reading through a book or tutorial on git, even if you don't remember the commands you learn, can help you learn some of the vocabulary³, and knowing that vocabulary will help you search for answers when you need them.

Finally, if you can find a way to do it, I think the best and easiest way to learn to use git and GitHub with R is to collaborate with someone who's used these tools before. Most of the time, these tools are very easy to use, but the small percent of the time that they're not, it can be significant stumbling blocks (in terms of the time it takes to figure out the fix) the first few times you use the tools, while someone familiar with them and working on the project can diagnose and get you over those bumps as you learn the ropes.

³ Some good git-related words to know to help you search for calls for rarer tasks: "commit", "branch", "merge", "revert", "push", "pull", "merge conflict", "remote", "origin", "master", "fork", "clone", "pull request".

5

Collect

For Environmental Health researchers, I think one of the most exciting developments in R recently is how it is changing how we can collect data, both for exposures and outcomes. One direction for this development is how researchers can collect and measure original data from experiments, including through new measurement technologies (e.g., phone-based Apps) and through new or rapidly changing health-related measurements (e.g., metabolomics, flow cytometry) and associated open-source software.

R is also facilitating and leveraging rapid developments in how researchers can access and query secondary data, including from large administrative databases, like those maintained by the US Census, NOAA, and the USGS. This section will provide an introduction to some of the ideas and techniques behind these developments for collecting secondary or public-use data for environmental health research, as well as give you some directions on where to go to find R packages that facilitate collecting open data from R.

5.1 Data packages

With R, you can create packages that are exclusively or mainly created to share data. Some researchers are creating and publishing this type of “data package”, and some may be relevant to your research. These packages provide access to open datasets that the package maintainer collected and processed and is now making available as an R package, but the data is static (as of the last version of the package), rather than interfacing to pull the most recent data. However, this set-up can achieve more stability—it won’t be broken by a change in an online API.

This package is too big for CRAN to host, so we have it posted through our own package repository.¹ Because of this, you’ll need to take a few extra steps to download and install the package:²

```
library(drat)
addRepo("geanders")
```

¹ For more on how and why we did this, see an article we wrote about the process for **The R Journal**.

² If you don’t have the drat package, install it in the usual way with `install.packages("drat")`.

```
install.packages("hurricaneexposedata")
install.packages("hurricaneexposure")
```

Once you've installed both the hurricaneexposedata and hurricaneexposure packages, you can load them as usual to access both the data and some functions to work with them:

```
library(hurricaneexposedata)
library(hurricaneexposure)
```

The hurricaneexposure package has a series of functions that let you explore different exposures during storms. For example, to get the storms where either New York County, NY, or Suffolk County, MA, (which includes Boston) were exposed to tropical storm-level winds (17.5 m/s or higher), you can run:³

```
knitr::include_graphics("images/tab_completion_example.png")

county_wind(counties = c("36061", "25025"), start_year = 1988,
            end_year = 2015, wind_limit = 17.5)
```

```
##      storm_id  fips vmax_sust vmax_gust
## 1    Bob-1991 25025  26.46639  39.43492
## 2    Bob-1991 36061  18.19559  27.11142
## 3  Bertha-1996 25025  29.64453  44.17035
## 4  Bertha-1996 36061  28.95496  43.14289
## 5    Floyd-1999 25025  24.46946  36.45949
## 6    Floyd-1999 36061  20.50178  30.54765
## 7    Hanna-2008 25025  18.36505  27.36392
## 8    Hanna-2008 36061  19.25390  28.68832
## 9    Irene-2011 36061  25.68553  38.27144
## 10 Sandy-2012 36061  21.99213  32.76827
##      sust_dur gust_dur closest_time_utc
## 1         210      570 1991-08-19 20:00
## 2          0      480 1991-08-19 15:00
## 3         240      525 1996-07-14 01:15
## 4         180      540 1996-07-13 19:45
## 5         345      750 1999-09-17 07:45
## 6          60      315 1999-09-17 00:15
## 7          0      150 2008-09-07 07:15
## 8          0      195 2008-09-07 01:45
## 9         165      510 2011-08-28 13:15
## 10        225      795 2012-10-30 00:30
##      storm_dist      local_time closest_date
## 1  27.042565 1991-08-19 16:00 1991-08-19
```

³ For any of these functions, you can find out what parameters to include, and in what format, but opening the helpfile for the function. For example, once you've loaded the hurricaneexposure package, you can open the helpfile for the county_wind function by calling ?county_wind. Also helpful for navigating packages: take advantage of the package::function notation and RStudio's tab completion to look up the names of functions in a package. For example, type hurricaneexposure::. A pop-up window should show up with all the functions in the package (press Tab if the pop-up doesn't automatically open).

> hurricaneexp

```
## 2 161.571830 1991-08-19 11:00 1991-08-19
## 3 38.177990 1996-07-13 21:15 1996-07-13
## 4 16.966013 1996-07-13 15:45 1996-07-13
## 5 51.254726 1999-09-17 03:45 1999-09-17
## 6 45.408483 1999-09-16 20:15 1999-09-16
## 7 6.202866 2008-09-07 03:15 2008-09-07
## 8 29.916672 2008-09-06 21:45 2008-09-06
## 9 5.796733 2011-08-28 09:15 2011-08-28
## 10 158.040788 2012-10-29 20:30 2012-10-29
```

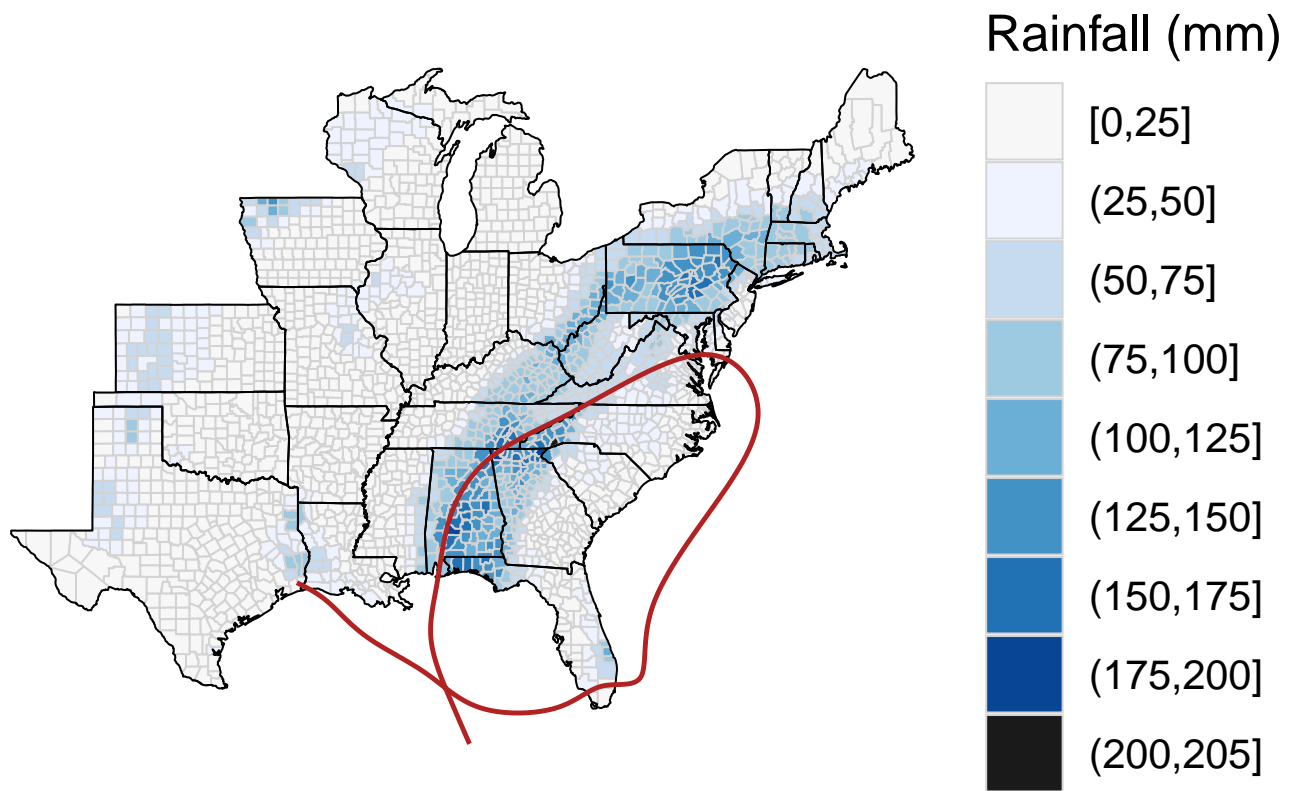
If you look up events based on flood events, you can instead run the `county_events` function:

```
county_events(counties = c("36061", "25025"),
              start_year = 1996, end_year = 2015, event_type = "flood")
```

```
##      fips      storm_id closest_time_utc
## 1 25025  Dennis-1999 1999-09-08 08:00
## 2 36061   Floyd-1999 1999-09-17 00:15
## 3 36061 Allison-2001 2001-06-17 14:15
## 4 36061 Frances-2004 2004-09-09 13:30
## 5 36061    Ivan-2004 2004-09-18 17:45
## 6 36061  Jeanne-2004 2004-09-29 06:30
## 7 36061   Beryl-2006 2006-07-20 22:15
## 8 36061   Barry-2007 2007-06-04 15:45
## 9 36061   Irene-2011 2011-08-28 13:15
## 10 25025   Sandy-2012 2012-10-29 22:00
## 11 25025  Andrea-2013 2013-06-08 11:30
## 12 36061  Andrea-2013 2013-06-08 06:45
##      storm_dist      local_time closest_date
## 1 390.047523 1999-09-08 04:00 1999-09-08
## 2 45.408483 1999-09-16 20:15 1999-09-16
## 3 158.909890 2001-06-17 10:15 2001-06-17
## 4 379.343696 2004-09-09 09:30 2004-09-09
## 5 311.346881 2004-09-18 13:45 2004-09-18
## 6 222.900157 2004-09-29 02:30 2004-09-29
## 7 207.358443 2006-07-20 18:15 2006-07-20
## 8 148.251718 2007-06-04 11:45 2007-06-04
## 9 5.796733 2011-08-28 09:15 2011-08-28
## 10 433.295980 2012-10-29 18:00 2012-10-29
## 11 45.412565 2013-06-08 07:30 2013-06-08
## 12 92.381282 2013-06-08 02:45 2013-06-08
```

The `hurricaneexposure` package also has functions for mapping the exposure data for specific storms. For example, to see the rainfall from Hurricane Ivan in 2004, you can run:

```
map_counties(storm = "Ivan-2004", metric = "rainfall")
```



From this map, you can see that the rain from this storm extended into New England, even after the storm looped back around to the east and south. This is why New York City had heavy rainfall (and flooding) from this event, but not tropical storm-level winds.

For more on the `hurricaneexposure` package, see its vignette.

5.2 Open data APIs

A range of environmental datasets are available online, especially through national agencies. For example, NOAA provides various weather datasets, while USGS has data on water quality,

You can visit webpages hosted by these agencies where you can download the datasets you need. However, this process can become tedious if you need lots of datasets, as may be the case for large studies incorporating many cities. Further, downloading the datasets “by hand” is hard to make reproducible, unless you meticulously write down all the steps you took as you visited the website. This means that your process will be harder for you to repeat in the future or for others to replicate.

A growing collection of R packages are now available that allow you to download datasets available online directly from R. This means that you can

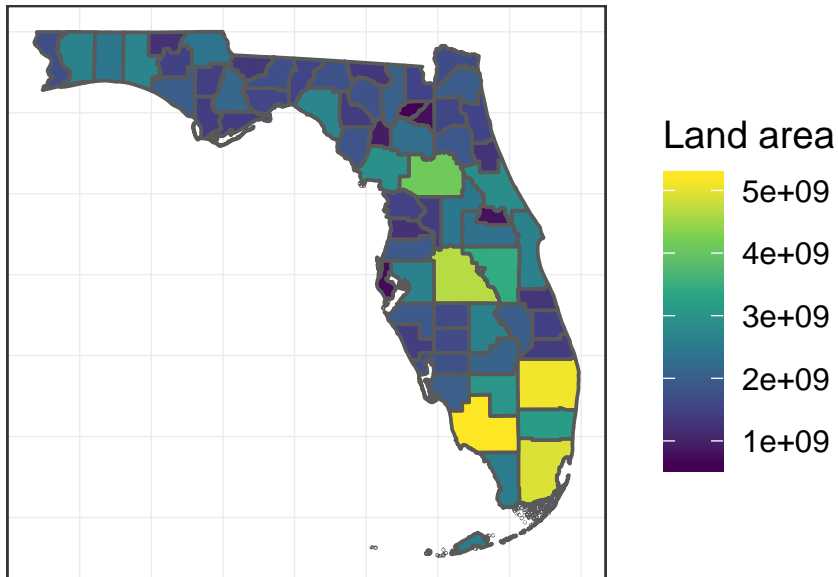
write an R script for your data collection, making this step both better documented and more reproducible.

One great example is the *tigris* package. This package lets you pull spatial data into R directly from the US Census. To get spatial data for all the counties in Florida, you can run:

```
library(tigris)
fl_counties <- counties(state = "FL", class = "sf")
```

Now you can plot this data with *ggplot*:

```
library(ggplot2)
ggplot() + geom_sf(data = fl_counties, aes(fill = ALAND)) +
  theme_bw() + scale_fill_viridis_c(name = "Land area",
    label = scales::scientific) + theme(axis.text = element_blank(),
    axis.ticks = element_blank())
```



5.3 Learn more

One of the best places to explore R packages for accessing open data for science is ROpenSci. Many of its packages facilitate access to databases of open data relevant to scientific research that have web services. You can browse through its packages on its Packages page. You may also want to check its affiliated Journal of Open-Source Software. Authors of these packages will also sometimes publish associated articles in The R Journal, so it's worth browsing through that occasionally. Finally, Twitter is a great place to keep an eye out for new packages, including those that can help you collect data. Follow the “#rstats” tag.

6

Bibliography

Bryan, J. (2018). Excuse me, do you have a moment to talk about version control? *The American Statistician*, 72(1):20–27.

Waller, L. A. (2018). Documenting and evaluating data science contributions in academic promotion in departments of statistics and biostatistics. *The American Statistician*, 72(1):11–19.

Wickham, H. (2015). *R packages: organize, test, document, and share your code.* ” O'Reilly Media, Inc.”.

Wickham, H. and Bryan, J. (2018). *usethis: Automate Package and Project Setup.* R package version 1.4.0.

Wickham, H., Hester, J., and Chang, W. (2019). *devtools: Tools to Make Developing R Packages Easier.* R package version 2.0.1.9000.